

✓ Ablation 1 : No self-training (evaluate just the initial weakly supervised model)

Final Project

Team 140: Anson Wong & Skyler MacAdam

GitHub link: <https://github.com/skylernovak/KeyClass/tree/main>

For our final project, we have selected the paper [Classifying Unstructured Clinical Notes via Automatic Weak Supervision](#). We will attempt to reproduce their study with similar results.

This is our first ablation from our reproduction of the KeyClass research paper. For additional context and information about the base approach and model, please see the primary notebook.

https://colab.research.google.com/drive/1ga1r8yUizxmBm_6Kpg1e-UyrPjQFXhix?usp=sharing

✓ Mount Notebook to Google Drive

This step needs to be completed the first time when Google Colab is opened and you are starting a new session.

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

✓ Methodology

Here, we begin our implementation of our study. Below you will find two sections, Data and Model.

Data

This section we will load and pre-process our data for our experiments.

Model

This section we will construct the model used for the experiment.

```
# Run this cell to install all requirements for this project.  
%cd /content/drive/My Drive/CS598FinalProject/  
%ls  
!pip install -r requirements.txt
```



```

Successfully uninstalled tokenizers-0.19.1
Attempting uninstall: transformers
Found existing installation: transformers 4.40.0
Uninstalling transformers-4.40.0:
Successfully uninstalled transformers-4.40.0
Successfully installed munkres-1.1.4 nvidia-cublas-cu12-12.1.3.1 nvidia-cuda-cupti

```

```

# Run this cell to import the remaining modules needed for this experiment.
%cd /content/drive/My Drive/CS598FinalProject/keyclass/
%ls

```

```

# Append paths for imports.
import sys
sys.path.append('/content/drive/My Drive/CS598FinalProject/keyclass/')
sys.path.append('/content/drive/My Drive/CS598FinalProject/scripts/')

```

```

# Import files and modules needed for KeyClass
import numpy as np
import utils
import models
import create_lfs
import train_classifier
import torch
import pickle
from os.path import join, exists
import pandas as pd
import os
import train_downstream_model

```

```

/content/drive/My Drive/CS598FinalProject/keyclass
create_lfs.py __init__.py models.py __pycache__/ train_classifier.py utils.py
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.

```

✓ Data

Source of the data

The data used in the paper is collected from 4 different datasets. Amazon, DBpedia, IMDB, and AG News. The authors collected this data and included a script with their paper to download copies of these datasets as well for others to reproduce the authors findings. We downloaded this data and included it within our google drive as well. These scripts are downloaded using the script `get_data.sh`.

These 4 datasets are used in common multiclass text classification problems, such as in movie review sentiments from IMBD.

Statistics

Utilizing the IMDB dataset, the paper attempts to classify movie reviews as either positive or negative. The KeyClass model was provided with common sense descriptions for these classes from industry experts. This constitutes minimal human input required for KeyClass to perform the classifications. For example, a positive review is typically associated with words such as "amazing", "exciting", or "fun" while negative reviews are more often associated with words such as "terrible", "boring", or "awful".

Both the train and test datasets contains 25000 sample movie reviews each. KeyClass attempts to label each review as either positive or negative based on these reviews. Each data split contains approximately 32 million characters, and the files are approximately 30 MB each.

Data process

The KeyClass framework consists of several key steps:

1. Finding relevant keywords and phrases from the class descriptions using a pre-trained language model (such as BERT).
2. Constructing labeling functions based on the identified keywords and using data programming to generate probabilistic labels for the training data.
3. Training a downstream classifier on the probabilistically labeled training data.
4. Self-training the downstream model on the entire training dataset to refine the classifier.

```
# Declare configuration variables
config_file_path = r'/content/drive/My Drive/CS598FinalProject/config_files/config_imdb.yml'
random_seed = 0 # Random seed for experiments
args = utils.Parser(config_file_path=config_file_path).parse()
```

The following block of code processes the text and creates embeddings. It is highly encouraged to use the Google Colab t4 GPU runtime type to execute this block of code. Using the T4 GPU, this takes a couple of minutes to execute. Using the CPU runtime type, this can take hours.

This block of code has been commented out as the embeddings have been saved in pickle files for later use in this notebook. This step is not required to be executed again, and is left here for informational / demonstrative purposes.

```
# args = utils.Parser(config_file_path=config_file_path).parse()

# if args['use_custom_encoder']:
#     model = models.CustomEncoder(pretrained_model_name_or_path=args['base_encoder'],
#                                   device='cuda' if torch.cuda.is_available() else 'cpu')
# else:
#     model = models.Encoder(model_name=args['base_encoder'],
#                             device='cuda' if torch.cuda.is_available() else 'cpu')

# for split in ['train', 'test']:
#     sentences = utils.fetch_data(dataset=args['dataset'], split=split, path=args['data_path'])
#     embeddings = model.encode(sentences=sentences, batch_size=args['end_model_batch_size'],
#                               show_progress_bar=args['show_progress_bar'],
#                               normalize_embeddings=args['normalize_embeddings'])
#     with open(join(args['data_path'], args['dataset'], f'{split}_embeddings.pkl'), 'wb') as f:
#         pickle.dump(embeddings, f)
```

Here we now load the training data, and then create the labeling function. Finally we create the probabilistic labels from the training document.

```

# Load training data
train_text = utils.fetch_data(dataset=args['dataset'], path=args['data_path'], split='train')

training_labels_present = False
if exists(join(args['data_path'], args['dataset'], 'train_labels.txt')):
    with open(join(args['data_path'], args['dataset'], 'train_labels.txt'), 'r') as f:
        y_train = f.readlines()
        y_train = np.array([int(i.replace('\n','')) for i in y_train])
        training_labels_present = True
else:
    y_train = None
    training_labels_present = False
    print('No training labels found!')

with open(join(args['data_path'], args['dataset'], 'train_embeddings.pkl'), 'rb') as f:
    X_train = pickle.load(f)

# Print dataset statistics
print(f"Getting labels for the {args['dataset']} data...")
print(f'Size of the data: {len(train_text)}')
if training_labels_present:
    print('Class distribution', np.unique(y_train, return_counts=True))

# Load label names/descriptions
label_names = []
for a in args:
    if 'target' in a: label_names.append(args[a])

# Creating labeling functions
labeler = create_lfs.CreateLabellingFunctions(base_encoder=args['base_encoder'],
                                              device=torch.device(args['device']),
                                              label_model=args['label_model'])
proba_preds = labeler.get_labels(text_corpus=train_text, label_names=label_names, min_df=args['min_df'],
                                ngram_range=args['ngram_range'], topk=args['topk'], y_train=y_train,
                                label_model_lr=args['label_model_lr'], label_model_n_epochs=args['label_model_n_epochs'],
                                verbose=True, n_classes=args['n_classes'])

y_train_pred = np.argmax(proba_preds, axis=1)

# Save the predictions
if not os.path.exists(args['preds_path']): os.makedirs(args['preds_path'])
with open(join(args['preds_path'], f"{args['label_model']}_proba_preds.pkl"), 'wb') as f:
    pickle.dump(proba_preds, f)

# Print statistics
print('Label Model Predictions: Unique value and counts', np.unique(y_train_pred, return_counts=True))
if training_labels_present:
    print('Label Model Training Accuracy', np.mean(y_train_pred==y_train))

# Log the metrics
training_metrics_with_gt = utils.compute_metrics(y_preds=y_train_pred, y_true=y_train, a

```

```
utils.log(metrics=training_metrics_with_gt, filename='label_model_with_ground_truth',  
          results_dir=args['results_path'], split='train')
```

excellently exquisite extremely well tabulous fairly good
 'fantastic' 'far best' 'film excellent' 'find good' 'fine job'
 'fine performances' 'finest' 'first rate' 'get good' 'give good'
 'gives best' 'gives good' 'gives great' 'good' 'good action' 'good also'
 'good although' 'good bad' 'good choice' 'good direction' 'good either'
 'good enough' 'good entertainment' 'good especially' 'good even'
 'good example' 'good film' 'good films' 'good first' 'good good'
 'good great' 'good idea' 'good movie' 'good music' 'good one' 'good ones'
 'good original' 'good part' 'good parts' 'good people' 'good performance'
 'good performances' 'good really' 'good reviews' 'good say' 'good show'
 'good special' 'good stuff' 'good taste' 'good thing' 'good things'
 'good think' 'good though' 'good tv' 'good use' 'good well' 'good work'
 'good would' 'good writing' 'got good' 'got great' 'great'
 'great character' 'great example' 'great film' 'great fun' 'great love'
 'great music' 'great one' 'great really' 'great show' 'great supporting'
 'great things' 'great time' 'greats' 'high quality' 'high rating'
 'highly recommend' 'highly recommended' 'however like' 'idea good'
 'like best' 'like good' 'like great' 'liked' 'liked one' 'looks great'
 'lot good' 'lot great' 'love good' 'lovely' 'luxury' 'made good'
 'made great' 'made well' 'make good' 'make great' 'makes good'
 'makes great' 'many good' 'many great' 'many reviewers' 'many reviews'
 'marvelously' 'may good' 'mean good' 'might good' 'movie excellent'
 'movie good' 'movie recommend' 'movie wonderful' 'much enjoyed'
 'much good' 'music good' 'music great' 'nearly good' 'nice' 'nice look'
 'nothing better' 'one best' 'one finest' 'one good' 'one great'
 'one like' 'overall good' 'overall think' 'particularly good'
 'people good' 'people like' 'performances good' 'perhaps best'
 'personal favorite' 'personally think' 'pleasant' 'positive reviews'
 'positive thing' 'possibly best' 'praise' 'prefer' 'prefers'
 'pretty decent' 'pretty good' 'probably best' 'probably good'
 'probably like' 'put good' 'qualities' 'quality' 'quality acting'
 'quite enjoyable' 'quite good' 'quite like' 'rather good' 'rating 10'
 'read review' 'read reviews' 'reading reviews' 'real good'
 'really appreciate' 'really enjoy' 'really enjoyed' 'really good'
 'really great' 'really like' 'really liked' 'really loved' 'really nice'
 'really recommend' 'recommend' 'recommend anyone' 'recommend everyone'
 'recommend film' 'recommend movie' 'recommend one' 'recommend see'
 'recommend watch' 'recommend watching' 'recommendation' 'recommended'
 'recommending' 'redeeming quality' 'reviews' 'satisfactory' 'say best'
 'say good' 'see good' 'seen good' 'show good' 'solid performances'
 'something better' 'something good' 'something interesting' 'splendid'
 'still enjoyable' 'still good' 'still great' 'strongly recommend'
 'surprisingly good' 'tasteful' 'terrific' 'thing good' 'think best'
 'think good' 'think great' 'though good' 'thought good' 'thought great'
 'time great' 'top notch' 'truly great' 'two best' 'want good'
 'watch good' 'well crafted' 'well good' 'well great' 'well made'
 'well produced' 'well worth' 'wonderful' 'wonderful film' 'wonderful job'
 'wonderful life' 'wonderful movie' 'wonderfully' 'worth look'
 'worth mentioning' 'worth seeing' 'worthwhile' 'would good'
 'would recommend']

==== Training the label model ====

100%|██████████| 100/100 [00:08<00:00, 11.33epoch/s]

Label Model Predictions: Unique value and counts (array([0, 1]), array([8914, 160

Label Model Training Accuracy 0.70016

Saving results in ../results/imdb/train_label_model_with_ground_truth_26-Apr-2024-

⌂ **B** **I** **<>** **🔗** **🖼️** **”** **☰** **☰** **—** **ψ** **😊** **📄**

```
### Train the Downstream Model
```

```
#### Find Class Descriptions
```

```
KeyClass starts with just the class descriptio
```

```
"negative words" without any labeled training
```

Train the Downstream Model

Find Class Descriptions

negative review) without any labeled training

Find Relevant Keywords

Using the class descriptions and a pre-trained KeyClass automatically extracts keywords and phrases of each class.

Probabilistically Label the Data

KeyClass uses the extracted keywords as labeling programming techniques to generate probabilistic dataset.

Train Downstream Model

KeyClass then trains a downstream text classification (e.g., a feed-forward neural network) using the probabilistically labeled training data. It initially only uses the most confidently labeled samples to train the model.

The key advantages of this approach are that it requires no manually labeled training data and the labeling functions are automatically generated in an interpretable way.

KeyClass starts with just the class descriptions (e.g., "positive review" and "negative review") without any labeled training data.

Find Relevant Keywords

Using the class descriptions and a pre-trained language model (like BERT), KeyClass automatically extracts keywords and phrases that are highly indicative of each class.

Probabilistically Label the Data

KeyClass uses the extracted keywords as labeling functions and applies data programming techniques to generate probabilistic labels for the entire training dataset.

Train Downstream Model

KeyClass then trains a downstream text classification model (e.g., a feed-forward neural network) using the probabilistically labeled training data. It initially only uses the most confidently labeled samples to train the model.

The key advantages of this approach are that it requires no manually labeled training data and the labeling functions are automatically generated in an interpretable way.

```

args = utils.Parser(config_file_path=config_file_path).parse()

# Set random seeds
random_seed = random_seed
torch.manual_seed(random_seed)
np.random.seed(random_seed)

X_train_embed_masked, y_train_lm_masked, y_train_masked, \
    X_test_embed, y_test, training_labels_present, \
    sample_weights_masked, proba_preds_masked = train_downstream_model.load_data(args)

# Train a downstream classifier

if args['use_custom_encoder']:
    encoder = models.CustomEncoder(pretrained_model_name_or_path=args['base_encoder'], device=device)
else:
    encoder = models.Encoder(model_name=args['base_encoder'], device=args['device'])

classifier = models.FeedForwardFlexible(encoder_model=encoder,
                                         h_sizes=args['h_sizes'],
                                         activation=eval(args['activation']),
                                         device=torch.device(args['device']))

print('\n==== Training the downstream classifier =====\n')
model = train_classifier.train(model=classifier,
                               device=torch.device(args['device']),
                               X_train=X_train_embed_masked,
                               y_train=y_train_lm_masked,
                               sample_weights=sample_weights_masked if args['use_noise_aware_loss'] else None,
                               epochs=args['end_model_epochs'],
                               batch_size=args['end_model_batch_size'],
                               criterion=eval(args['criterion']),
                               raw_text=False,
                               lr=eval(args['end_model_lr']),
                               weight_decay=eval(args['end_model_weight_decay']),
                               patience=args['end_model_patience'])

end_model_preds_train = model.predict_proba(torch.from_numpy(X_train_embed_masked), batch_size=batch_size)
end_model_preds_test = model.predict_proba(torch.from_numpy(X_test_embed), batch_size=512, raw_text=False)

Confidence of least confident data point of class 0: 0.9118951704039087
Confidence of least confident data point of class 1: 0.9999157389338196

==== Data statistics ====
Size of training data: (25000, 768), testing data: (25000, 768)
Size of testing labels: (25000,)
Size of training labels: (25000,)
Training class distribution (ground truth): [0.5 0.5]
Training class distribution (label model predictions): [0.35656 0.64344]

KeyClass only trains on the most confidently labeled data points! Applying mask...

```

```
==== Data statistics (after applying mask) ====
```

```
Size of training data: (7000, 768)
```

```
Size of training labels: (7000,)
```

```
Training class distribution (ground truth): [0.55057143 0.44942857]
```

```
Training class distribution (label model predictions): [0.5 0.5]
```

```
===== Training the downstream classifier =====
```

```
Epoch 18: 90%|██████████| 18/20 [00:05<00:00, 3.55batch/s, best_loss=0.546, running_lc
```

✓ ABLATION

No self-training (evaluate just the initial weakly supervised model)

To evaluate the performance of the initial weakly supervised model without self-training, I remove the code block that performs self-training.

```
# # Fetching the raw text data for self-training
# X_train_text = utils.fetch_data(dataset=args['dataset'], path=args['data_path'], split='tr
# X_test_text = utils.fetch_data(dataset=args['dataset'], path=args['data_path'], split='tes

# model = train_classifier.self_train(model=model,
#                                     X_train=X_train_text,
#                                     X_val=X_test_text,
#                                     y_val=y_test,
#                                     device=torch.device(args['device']),
#                                     lr=eval(args['self_train_lr']),
#                                     weight_decay=eval(args['self_train_weight_decay']),
#                                     patience=args['self_train_patience'],
#                                     batch_size=args['self_train_batch_size'],
#                                     q_update_interval=args['q_update_interval'],
#                                     self_train_thresh=eval(args['self_train_thresh']),
#                                     print_eval=True)

# end_model_preds_test = model.predict_proba(X_test_text, batch_size=args['self_train_batch_

# # Print statistics
# testing_metrics = utils.compute_metrics_bootstrap(y_preds=np.argmax(end_model_preds_test,
#                               y_true=y_test,
#                               average=args['average'],
#                               n_bootstrap=args['n_bootstrap'],
#                               n_jobs=args['n_jobs'])
# print(testing_metrics)
```

✓ After training the downstream classifier and (and skipping self-training), the notebook performs the following evaluation steps:

1. Load Test Data: It loads the test dataset, including the test text samples (`X_test_embed`) and the ground truth test labels (`y_test`).
2. Evaluate Trained Model on Test Set: It uses the trained downstream model to make predictions on the test set, obtaining the test set predictions (`end_model_preds_test`).
3. Compute Test Metrics: It computes various performance metrics on the test set predictions, including:

Metrics using ground truth labels (`y_test`): Compute metrics using `utils.compute_metrics_bootstrap()`, which does bootstrap sampling to get confidence intervals. This provides an assessment of the model's true performance on the test set.

Metrics using label model predictions (`y_train_lm_masked`): Compute metrics using `utils.compute_metrics()`. This shows how the model performs compared to the noisy labels used for training.

Print Test Metrics: The notebook prints out the test set performance metrics, showing the model's accuracy, precision, recall, and F1-score.

Self-Train the Model: Finally, it loads the self-trained model checkpoint and evaluates the self-trained model on the test set, printing the updated test set performance metrics.

```

# end_model_path='/content/drive/My Drive/CS598FinalProject/models/end_model.pth'
# end_model_self_trained_path='/content/drive/My Drive/CS598FinalProject/models/end_model_se

args = utils.Parser(config_file_path=config_file_path).parse()

# Set random seeds
random_seed = random_seed
torch.manual_seed(random_seed)
np.random.seed(random_seed)

X_train_embed_masked, y_train_lm_masked, y_train_masked, \
    X_test_embed, y_test, training_labels_present, \
    sample_weights_masked, proba_preds_masked = train_downstream_model.load_data(args)

# model = torch.load(end_model_path)

end_model_preds_train_key_class = model.predict_proba(torch.from_numpy(X_train_embed_masked))
end_model_preds_test_key_class = model.predict_proba(torch.from_numpy(X_test_embed), batch_s

# Print statistics
if training_labels_present:
    training_metrics_with_gt = utils.compute_metrics(y_preds=np.argmax(end_model_preds_train_key
                                                    y_true=y_train_masked,
                                                    average=args['average'])
    print('training_metrics_with_gt', training_metrics_with_gt)

training_metrics_with_lm = utils.compute_metrics(y_preds=np.argmax(end_model_preds_train_key
                                                    y_true=y_train_lm_masked,
                                                    average=args['average'])
print('training_metrics_with_lm', training_metrics_with_lm)

testing_metrics = utils.compute_metrics_bootstrap(y_preds=np.argmax(end_model_preds_test_key
                                                    y_true=y_test,
                                                    average=args['average'],
                                                    n_bootstrap=args['n_bootstrap'],
                                                    n_jobs=args['n_jobs'])

print('testing_metrics', testing_metrics)

print('\n==== Self-training the downstream classifier =====\n')

# Fetching the raw text data for self-training
X_train_text = utils.fetch_data(dataset=args['dataset'], path=args['data_path'], split='train')
X_test_text = utils.fetch_data(dataset=args['dataset'], path=args['data_path'], split='test')

# model = torch.load(end_model_self_trained_path)

end_model_preds_test_self_training = model.predict_proba(X_test_text, batch_size=args['self_

# Print statistics

```

```

testing_metrics = utils.compute_metrics_bootstrap(y_preds=np.argmax(end_model_preds_test_sel
                                                    y_true=y_test,
                                                    average=args['average'],
                                                    n_bootstrap=args['n_bootstrap'],
                                                    n_jobs=args['n_jobs'])
print('testing_metrics after self train', testing_metrics)

```

Confidence of least confident data point of class 0: 0.9118951704039087
 Confidence of least confident data point of class 1: 0.9999157389338196

==== Data statistics ====

Size of training data: (25000, 768), testing data: (25000, 768)

Size of testing labels: (25000,)

Size of training labels: (25000,)

Training class distribution (ground truth): [0.5 0.5]

Training class distribution (label model predictions): [0.35656 0.64344]

KeyClass only trains on the most confidently labeled data points! Applying mask...

==== Data statistics (after applying mask) ====

Size of training data: (7000, 768)

Size of training labels: (7000,)

Training class distribution (ground truth): [0.55057143 0.44942857]

Training class distribution (label model predictions): [0.5 0.5]

training_metrics_with_gt [0.9204285714285714, 0.9241398875739136, 0.9204285714285714]

training_metrics_with_lm [0.9218571428571428, 0.9219690593856281, 0.9218571428571428]

[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n_jobs=10)]: Done 40 tasks | elapsed: 0.9s

[Parallel(n_jobs=10)]: Done 100 out of 100 | elapsed: 1.8s finished

testing_metrics [[0.8491612 0.00204014]

[0.86128409 0.00184495]

[0.8491612 0.00204014]]

===== Self-training the downstream classifier =====

[Parallel(n_jobs=10)]: Using backend LokyBackend with 10 concurrent workers.

[Parallel(n_jobs=10)]: Done 30 tasks | elapsed: 17.0s

testing_metrics after self train [[0.8495204 0.00235398]

[0.8615633 0.00208607]

[0.8495204 0.00235398]]

[Parallel(n_jobs=10)]: Done 100 out of 100 | elapsed: 20.5s finished

✓ Save Ablations Result

No self-training (evaluate just the initial weakly supervised model)


```

# File paths for saving the arrays
save_paths = {
    'y_test_ablation_1': '/content/drive/My Drive/CS598FinalProject/y_test_ablation_1.npy',
    'y_train_masked_ablation_1': '/content/drive/My Drive/CS598FinalProject/y_train_masked_a',
    'y_train_lm_masked_ablation_1': '/content/drive/My Drive/CS598FinalProject/y_train_lm_ma',
    'end_model_preds_train_key_class_ablation_1': '/content/drive/My Drive/CS598FinalProject',
    'end_model_preds_test_key_class_ablation_1': '/content/drive/My Drive/CS598FinalProject/',
    'end_model_preds_test_self_training_ablation_1': '/content/drive/My Drive/CS598FinalProj
}

# Arrays to save
arrays_to_save = {
    'y_test_ablation_1': y_test,
    'y_train_masked_ablation_1': y_train_masked,
    'y_train_lm_masked_ablation_1': y_train_lm_masked,
    'end_model_preds_train_key_class_ablation_1': end_model_preds_train_key_class,
    'end_model_preds_test_key_class_ablation_1': end_model_preds_test_key_class,
    'end_model_preds_test_self_training_ablation_1': end_model_preds_test_self_training
}

# Save arrays
for name, array in arrays_to_save.items():
    np.save(save_paths[name], array)

```

✓ Load Ablation Results

```

file_path = '/content/drive/My Drive/CS598FinalProject/y_test_ablation_1.npy'
y_test_ablation_1 = np.load(file_path)

```

```

file_path = '/content/drive/My Drive/CS598FinalProject/y_train_masked_ablation_1.npy'
y_train_masked_ablation_1 = np.load(file_path)

```

```

file_path = '/content/drive/My Drive/CS598FinalProject/y_train_lm_masked_ablation_1.npy'
y_train_lm_masked_ablation_1 = np.load(file_path)

```

```

file_path = '/content/drive/My Drive/CS598FinalProject/end_model_preds_train_key_class_ablat
end_model_preds_train_key_class_ablation_1 = np.load(file_path)

```

```

file_path = '/content/drive/My Drive/CS598FinalProject/end_model_preds_test_key_class_ablati
end_model_preds_test_key_class_ablation_1 = np.load(file_path)

```

```

file_path = '/content/drive/My Drive/CS598FinalProject/end_model_preds_test_self_training_at
end_model_preds_test_self_training_ablation_1 = np.load(file_path)

```

```
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
# weak supervision sources (keywords and phrases) results
```

```
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_curve, auc
import seaborn as sns

# Plot confusion matrix
def plot_confusion_matrix(y_true, y_pred, title='Confusion Matrix', labels=None):
    cm = confusion_matrix(y_true, y_pred)
    if labels:
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
    else:
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(title)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.show()

# Plot precision-recall curve
def plot_precision_recall_curve(y_true, y_score, title):
    precision, recall, _ = precision_recall_curve(y_true, y_score)
    plt.plot(recall, precision, marker='.')
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title(title)
    plt.show()

# Plot ROC curve
def plot_roc_curve(y_true, y_score, title):
    fpr, tpr, _ = roc_curve(y_true, y_score)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label='ROC curve (AUC = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Random')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(title)
    plt.legend(loc='lower right')
    plt.show()

# Plot histogram of predicted probabilities
def plot_predicted_probabilities_histogram(y_prob):
    plt.hist(y_prob, bins=10)
    plt.xlabel('Predicted Probability')
    plt.ylabel('Frequency')
    plt.title('Histogram of Predicted Probabilities')
    plt.show()

# Example usage
y_preds_ablation_1=np.argmax(end_model_preds_train_key_class_ablation_1, axis=1)

# plot_predicted_probabilities_histogram(end_model_preds_train_key_class)
```